

Technical Reference

```
{r setup, include = FALSE} knitr::opts_chunk$set(echo = TRUE,
tidy.opts = list(width.cutoff = 80), tidy = TRUE) # Outline
```

1. [About](#)
 - a. [Getting started](#)
 - b. [Output abbreviations](#)
 - c. [Project structure](#)
2. [Setup](#)
 - a. [Dependencies](#)
 - b. [Fields](#)
 - c. [Year](#)
 - d. [States](#)
 - e. [Counties](#)
 - f. [Census API key](#)
 - g. [Functions](#)
 1. [Override base and stats function defaults](#)
 2. [Create custom half-standard deviation breaks](#)
 3. [Exception](#)
 4. [Move column or vector of columns to last position](#)
 5. [Summarize data](#)
3. [Variance replicate table download](#)
 - a. [Download variance replicates from Census website](#)
 - b. [Combine and format downloads](#)
4. [Variance replicate table processing](#)
 - a. [Compute racial minority count MOE](#)
 - b. [Save results](#)
5. [ACS estimates download](#)
 - a. [Fields](#)
 - b. [Download counts and universes from Census API](#)
 1. [Exception](#)
 - c. [Download percentages from Census API](#)
 - d. [Format downloads](#)
 1. [Exception](#)
 2. [Exception](#)
 3. [Exception](#)
6. [ACS estimates calculations](#)
 - a. [Percentages and percentage MOEs](#)
 1. [Calculation](#)

- 2. Result
 - 3. *Exception*
 - 4. *Exception*
- b. Percentile
 - 1. Calculation
 - 2. Result
- c. IPD score and classification
 - 1. Calculation
 - 2. Result
- d. Composite IPD score
 - 1. Calculation
 - 2. Result
- 7. ACS estimates cleaning
- 8. Summary tables
 - a. Counts by indicator
 - b. Breaks by indicator
 - c. Summary by indicator
 - d. County means by indicator
- 9. Export
 - a. Append to TIGER/LINE file
 - b. Export files

1. About

DVRPC's IPD analysis identifies populations of interest under Title VI of the Civil Rights Act and the Executive Order on Environmental Justice (#12898) using 2013-2017 American Community Survey (ACS) five-year estimates from the U.S. Census Bureau. IPD analysis assists both DVRPC and outside organizations in equity work by identifying populations of interest, including youth, older adults, female, racial minority, ethnic minority, foreign-born, limited English proficiency, disabled, and low-income populations at the census tract level in DVRPC's nine-county region.

There are many ways of identifying these populations of interest. This document discusses DVRPC's process, which is automated in an R script.

1a. Getting started

For guidance on software prerequisites and how to run this script, see `getting_started.pdf` in the documentation folder.

1b. Output abbreviations

Components of field names that you'll see in outputs and throughout the script.

```
{r table1, tidy=FALSE} library(tidyverse) data1 <-
tribble( ~"Component", ~"Equivalent", "D", "Disabled", "EM",
"Ethnic Minority", "F", "Female", "FB", "Foreign-
Born", "LEP", "Limited English Proficiency", "LI",
"Low-Income", "OA", "Older Adults", "RM", "Racial
Minority", "Y", "Youth", "CntEst", "Count Estimate",
"CntMOE", "Count MOE", "PctEst", "Percentage Estimate",
"PctMOE", "Percentage MOE", "Pctile", "Percentile", "Score",
"Score", "Class", "Classification" ) knitr::kable((data1),
booktabs = TRUE, caption = 'Component and Equivalent')
```

Abbreviations of field names that you'll see in outputs *not* comprised of the above components.

```
{r table2, tidy=FALSE} library(tidyverse) data2 <-
tribble( ~"Abbreviation", ~"Equivalent", "GEOID", "Census
Tract Identifier", "STATEFP", "State FIPS Code", "COUNTYFP",
"County FIPS Code", "NAME", "Census Tract FIPS Code",
"IPD_Score", "Composite IPD Score", "U_TPopEst", "Total
Population Estimate", "U_TPopMOE", "Total Population MOE",
"U_Pop5Est", "Population 5+ Estimate", "U_Pop5MOE",
"Population 5+ MOE", "U_PPovEst", "Poverty Status Population
Estimate", "U_PPovMOE", "Poverty Status Population MOE",
"U_PNICEst", "Non-Institutional Civilian Population Estimate",
"U_PNICMOE", "Non-Institutional Civilian Population MOE", )
knitr::kable((data2), booktabs = TRUE, caption = 'Abbreviation and
Equivalent')
```

1c. Project structure

This script uses relative file paths based off the location of `ipd.Rproj`. As long as you download the entire repository, the script should have no trouble locating the correct subfolders. All of the subsequent years files are based on the same architecture. The project is structured as follows:

```
{r file_structure, eval = FALSE} ipd ipd.Rproj script.R
documentation discussion.pdf getting_started.pdf
script_reference.pdf script_reference.Rmd variables.csv
outputs breaks_by_indicator.csv counts_by_indicator.csv
ipd.csv ipd.dbf ipd.prj ipd.shp ipd.shx
mean_by_county.csv summary_by_indicator.csv
```

2. Setup

2a. Dependencies

Packages required to run this script. If you don't have the packages, you'll get the warning
 Error in library (<name of package>) : there is no package called
 '<name of package>', in which case you'll need to install the package before proceeding.

```
{r packages, message = FALSE} library(plyr); library(here);
library(sf); library(summarytools); library(tidycensus);
library(tidyverse); library(tigris); library(dplyr); library(descr);
```

2b. Fields

The base information we need for IPD analysis are universes, counts, and percentages for nine indicators at the census tract level. For each indicator, the table below shows the indicator name, its abbreviation used in the script, its universe, its count, and its percentage field if applicable. Because the schemata of ACS tables can change with each annual ACS update, these field names are applicable *only* to 2013-2017 ACS Five-Year Estimates.

Some percentage fields are empty. This is okay: we will compute the percentages when they are not directly available from the ACS.

Note that variable B02001_002 (“Estimate; Total: - White alone”) is listed as the count for Racial Minority. This is a mathematical shortcut: otherwise, we would need to add several subfields to compute the same estimate. The desired count is B02001_001 (Universe) – B02001_002 (“Estimate; Total: - White alone”). The subtraction is computed after download in Section 5d.i., making a correct estimate and an incorrect MOE. The correct MOE for the count, as calculated in Section 4, will be appended later.

```
{r table1, tidy=FALSE} library(tidyverse) data3 <-
tribble( ~"Indicator", ~"Abbreviation", ~"Universe", ~"Count",
~"Percentage", "Disabled", "D", "S1810_C01_001", "S1810_C02_001",
"S1810_C03_001", "Ethnic Minority", "EM", "B03002_001", "B03002_012",
"N/A", "Female", "F", "S0101_C01_001", "S0101_C05_001",
"DP05_0003PE", "Foreign-Born", "FB", "B05012_001", "B05012_003",
"N/A", "Limited English Proficiency", "LEP", "S1601_C01_001",
"S1601_C05_001", "S1601_C06_001", "Low-Income", "LI", "S1701_C01_001",
"S1701_C01_042", "N/A", "Older Adults", "OA", "S0101_C01_001",
"S0101_C01_030", "S0101_C02_030", "Racial Minority", "RM",
"B02001_001", "B02001_002", "N/A", "Youth", "Y", "B03002_001",
"B09001_001", "N/A" ) knitr::kable((data3), booktabs = TRUE, caption =
'')
```

While it’s quicker to embed the names of the desired columns into the code, fields are explicitly spelled out in this script. This is a purposeful design choice. The user should check that the field names point to the correct API request with every IPD update. The best way to check the field names is to visit Census Developers ([link](#)) and select the corresponding API. For a history of the ACS variables used in IPD 2015, 2016, and 2017, see `variables.csv` in the documentation folder.

```
disabled_universe <- “S1810_C01_001”
```

```
disabled_count <- “S1810_C02_001”
```

```
disabled_percent <- “S1810_C03_001”
```

```
ethnic_minority_universe <- "B03002_001"
ethnic_minority_count <- "B03002_012"
ethnic_minority_percent <- NA
female_universe <- "S0101_C01_001"
female_count <- "S0101_C05_001"
female_percent <- "DP05_0003PE"
foreign_born_universe <- "B05012_001"
foreign_born_count <- "B05012_003"
foreign_born_percent <- NA
limited_english_proficiency_universe <- "S1601_C01_001"
limited_english_proficiency_count <- "S1601_C05_001"
limited_english_proficiency_percent <- "S1601_C06_001"
low_income_universe <- "S1701_C01_001"
low_income_count <- "S1701_C01_042"
low_income_percent <- NA
older_adults_universe <- "S0101_C01_001"
older_adults_count <- "S0101_C01_030"
older_adults_percent <- "S0101_C02_030"
racial_minority_universe <- "B02001_001"
racial_minority_count <- "B02001_002"
racial_minority_percent <- NA
youth_universe <- "B03002_001"
youth_count <- "B09001_001"
youth_percent <- NA
```

2c. Year

```
The data download year. {r year} ipd_year <- 2017
```

2d. States

The data download state or states. Use the two-character text abbreviation. {r states}
`ipd_states <- c("NJ", "PA")`

2e. Counties

The counties in your study area. Use five-digit characters concatenating the two-digit state and three-digit county FIPS codes. {r counties} `ipd_counties <- c("34005", "34007", "34015", "34021", "42017", "42029", "42045", "42091", "42101")`

2f. Census API Key

Placeholder if you have never installed an API key before. If this is your first time accessing the Census API using R, see `getting_started.pdf` in the documentation folder. {r api_key} # Census API Key # `census_api_key("YOUR API KEY GOES HERE", install = TRUE)`

THE TYPICAL USER SHOULD NOT HAVE TO EDIT ANYTHING BELOW THIS POINT.

2g. Functions

Load custom functions.

2g.i. Override base and stats function defaults

A time-saver so that it's not required to call `na.rm = TRUE` every time common functions are called. {r override} `min <- function(i, ..., na.rm = TRUE) { base::min(i, ..., na.rm = na.rm) } mean <- function(i, ..., na.rm = TRUE) { base::mean(i, ..., na.rm = na.rm) } sd <- function(i, ..., na.rm = TRUE) { stats::sd(i, ..., na.rm = na.rm) } max <- function(i, ..., na.rm = TRUE) { base::max(i, ..., na.rm = na.rm) }`

2g.ii. Create custom half-standard deviation breaks

For a given vector of numbers `x` and a number of bins `i`, `st_dev_breaks` computes the bin breaks starting at $-0.5 \cdot stdev$ and $0.5 \cdot stdev$. For the purposes of IPD analysis, `i = 5`, and `st_dev_breaks` calculates the minimum, $-1.5 \cdot stdev$, $-0.5 \cdot stdev$, $0.5 \cdot stdev$, $1.5 \cdot stdev$, and maximum values. These values are later used to slice the vector into five bins.

2g.iii. Exception

All minima are coerced to equal zero. If the first bin break ($-1.5 \cdot stdev$) is negative, as happens when the data has a large spread and therefore a large standard deviation, then this bin break is coerced to equal 0.1. In these cases, only estimates of 0 percent will be

```
placed in the bottom bin. {r st_dev_breaks} st_dev_breaks <- function(x, i,
na.rm = TRUE){ half_st_dev_count <- c(-1 * rev(seq(1, i, by = 2)),
seq(1, i, by = 2)) if((i %% 2) == 1) { half_st_dev_breaks <-
sapply(half_st_dev_count, function(i)
(0.5 * i * sd(x)) + mean(x)) half_st_dev_breaks[[1]] <- 0
half_st_dev_breaks[[2]] <- ifelse(half_st_dev_breaks[[2]] < 0,
0.1, half_st_dev_breaks[[2]])
half_st_dev_breaks[[i + 1]] <- ifelse(max(x) > half_st_dev_breaks[[i +
1]], max(x),
half_st_dev_breaks[[i + 1]]) } else { half_st_dev_breaks <- NA
} return(half_st_dev_breaks) }
```

2g.iv. Move column or vector of columns to last position

The requested schema for IPD data export renames and places all relevant universes in the final columns of the dataset. `move_last` moves a column or vector of column names to the last position(s) in a data frame. {r move_last} `move_last <- function(df, last_col) { match(c(setdiff(names(df), last_col), last_col), names(df)) }`

2g.v. Summarize data

`description` tailors the exports from `summarytools::descr` to create summary tables with the requested fields. `0.5 * stdev` is returned after `stdev`. {r description} `description <- function(i) { des <- as.numeric(descr(i, na.rm = TRUE, stats = c("min", "med", "mean", "sd", "max"))) des <- c(des[1:4], des[4] / 2, des[5]) return(des) }`

3. Variance replicate table download

This will feel out of order, but it's necessary. The racial minority indicator is created by summing up several subgroups in ACS Table B03002. This means that the MOE for the count has to be computed. While the ACS has issued guidance on computing the MOE by aggregating subgroups, using the approximation formula can artificially deflate the derived MOE. Variance replicate tables are used instead to account for covariance and compute a more accurate MOE. The MOE computed from variance replicates is substituted in for the racial minority count MOE in Section 5d.ii.

See the Census Bureau's Variance Replicate Tables Documentation ([link](#)) for additional guidance on working with variance replicates.

3a. Download variance replicates from Census website

Download, unzip, and read variance replicate tables for Table B02001. Results are combined into a single table called `var_rep`. {r varrep_download, tidy = TRUE, message = FALSE} `ipd_states_numeric <- fips_codes %>% filter(state %in% ipd_states) %>% select(state_code) %>% distinct(.) %>% pull(.)` `var_rep <- NULL for (i in 1:length(ipd_states)){ url <- paste0("https://www2.census.gov/programs-surveys/acs/replicate_estimat`

```
es/",
      ipd_year,
      "/data/5-year/140/B02001_",
      ipd_states_numeric[i],
      ".csv.gz") temp <- tempfile() download.file(url, temp) var_rep_i
<- read_csv(gzfile(temp)) var_rep <- rbind(var_rep, var_rep_i) }
```

3b. Combine and format downloads

Subset `var_rep` for the study area defined in `ipd_counties` and extract the necessary subgroups. {r varrep_merge, message = FALSE} var_rep <- var_rep %>%
mutate_at(vars(GEOID), funs(str_sub(., 8, 18))) %>%
filter(str_sub(GEOID, 1, 5) %in% ipd_counties) %>% select(-TBLID, -
NAME, -ORDER, -moe, -CME, -SE) %>% filter(TITLE %in% c("Black or
African American alone", "American Indian and
Alaska Native alone", "Asian alone",
"Native Hawaiian and Other Pacific Islander alone",
"Some other race alone", "Two or more races:"))

4. Variance replicate table processing

4a. Compute racial minority count MOE

Add up the racial minority counts into a single count per census tract for the estimate and 80 variance replicates. Separate the resulting data frame into estimates and variance replicates. {r varrep_subset, message = FALSE} num <- var_rep %>%
group_by(GEOID) %>% summarize_if(is.numeric, funs(sum)) %>%
select(-GEOID) estim <- num %>% select(estimate) individual_replicate
<- num %>% select(-estimate) Compute the variance replicate for the count. GEOIDs
are stored as `id` to be re-appended to the MOEs after they are calculated. {r
varrep_calc, message = FALSE} id <- var_rep %>% select(GEOID) %>%
distinct(.) %>% pull(.) sqdiff_fun <- function(v, e) (v - e) ^ 2
sqdiff <- mapply(sqdiff_fun, individual_replicate, estim) sum_sqdiff
<- rowSums(sqdiff) variance <- 0.05 * sum_sqdiff moe <-
round(sqrt(variance) * 1.645, 0)

4b. Save results

Save the racial minority MOE. {r varrep_save, message = FALSE} rm_moe <-
cbind(id, moe) %>% as_tibble(.) %>% rename(GEOID10 = id, RM_CntMOE
= moe) %>% mutate_at(vars(RM_CntMOE), as.numeric) Here are the first few
lines of `rm_moe`: {r varrep_preview} head(rm_moe)

5. ACS estimates download

5a. Fields

Fields for downloads from the ACS API were discussed in Section 2b.

5b. Download counts and universes from Census API

Download counts and percentages for each of IPD's nine indicators. Note that the download is for all census tracts in `ipd_states`.

Input data for IPD comes from ACS Subject Tables, Detailed Tables, and Data Profiles. While one can request all the fields for Subject Tables in one batch, mixing requests for two or more different types of tables will result in failure. For this reason, the counts and universe fields supplied by the user in Section 2b are evaluated for their contents and split into three batches: `s_counts` for Subject Tables, `d_counts` for Detailed Tables, and `dp_counts` for Data Profiles.

The chunk below zips the user-defined calls from the API with the output abbreviations into a data frame called `counts_calls` and separates the calls into three batches. {r
`api_counts, message = FALSE} counts <- c(disabled_count,
disabled_universe, ethnic_minority_count,
ethnic_minority_universe, female_count, female_universe,
foreign_born_count, foreign_born_universe,
limited_english_proficiency_count,
limited_english_proficiency_universe, low_income_count,
low_income_universe, older_adults_count,
older_adults_universe, racial_minority_count,
racial_minority_universe, youth_count, youth_universe)
counts_ids <- c("D_C", "D_U", "EM_C", "EM_U", "F_C", "F_U",
"FB_C", "FB_U", "LEP_C", "LEP_U", "LI_C", "LI_U",
"OA_C", "OA_U", "RM_C", "RM_U", "Y_C", "Y_U") counts_calls <-
tibble(id = counts_ids, api = counts) %>% drop_na(.) s_calls <-
counts_calls %>% filter(str_sub(api, 1, 1) == "S") d_calls <-
counts_calls %>% filter(str_sub(api, 1, 1) == "B") dp_calls <-
counts_calls %>% filter(str_sub(api, 1, 1) == "D") API calls are made
separately for ACS Subject Tables, Detailed Tables, and Data Profiles and appended to
dl_counts. Sometimes there are no requests for an ACS table type; in these situations, the
script bypasses a download attempt. Then, information from counts_calls is used to
rename the downloads to the appropriate abbreviation. {r api_counts_calls,
message = FALSE} dl_counts <- NULL if(length(s_calls$id > 0))
{ s_counts <- get_acs(geography = "tract",
state = ipd_states, output = "wide",
year = ipd_year, variables = s_calls$api) %>%
select(-NAME) dl_counts <- bind_cols(dl_counts, s_counts) }
if(length(d_calls$id > 0)){ d_counts <- get_acs(geography = "tract",
state = ipd_states, output = "wide",
year = ipd_year, variables = d_calls$api) %>%
select(-NAME) dl_counts <- left_join(dl_counts, d_counts) }
if(length(dp_calls$id > 0)){ dp_counts <- get_acs(geography =
"tract", state = ipd_states,
output = "wide", year = ipd_year,
variables = dp_calls$api) %>% select(-NAME) dl_counts <-
left_join(dl_counts, dp_counts) } counts_calls$api <-
str_replace(counts_calls$api, "E$", "") for(i in`

```

1:length(counts_calls$id)){ names(dl_counts) <-
str_replace(names(dl_counts),
counts_calls$api[i],
counts_calls$id[i]) } dl_counts <- dl_counts %>% rename(GE0ID10 =
GE0ID)

```

5b.i. Exception

The API does not allow redundant downloads, so universes for Older Adults and Youth are duplicated after download. `duplicate_cols` identifies duplicate API calls, and `combined_rows` serves as a crosswalk to duplicate and rename fields. {r
`api_counts_duplicator`} `duplicate_cols <- counts_calls %>%`
`group_by(api) %>% filter(n()>1) %>% summarize(orig = id[1],`
`duplicator = id[2]) e_paste <- function(i) paste0(i, "E") m_paste <-`
`function(i) paste0(i, "M") e_rows <- apply(duplicate_cols, 2, e_paste)`
`m_rows <- apply(duplicate_cols, 2, m_paste) combined_rows <-`
`as_tibble(rbind(e_rows, m_rows)) %>% mutate_all(as.character) for(i`
`in 1:length(combined_rows$api))`
`{ dl_counts[combined_rows$duplicator[i]] <-`
`dl_counts[combined_rows$orig[i]] }`

5c. Download percentages from Census API

Download percentage tables that are available for four of IPD's nine indicators. We will compute percentages and their associated MOEs for the rest of the dataset later. The procedure is identical to that described in Section 5b. {r `api_percs`, `message = FALSE`} `perc_s <- c(disabled_percent,`
`ethnic_minority_percent,`
`foreign_born_percent,`
`low_income_percent,`
`racial_minority_percent,`
`c("D_P", "EM_P", "F_P", "FB_P", "LEP_P",`
`"OA_P", "RM_P", "Y_P") perc_s_calls <- tibble(id = perc_s_ids, api =`
`perc_s) %>% drop_na(.) s_calls <- perc_s_calls %>%`
`filter(str_sub(api, 1, 1) == "S") d_calls <- perc_s_calls %>%`
`filter(str_sub(api, 1, 1) == "B") dp_calls <- perc_s_calls %>%`
`filter(str_sub(api, 1, 1) == "D") dl_percs <- NULL`
`if(length(s_calls$id > 0)){ s_percs <- get_acs(geography = "tract",`
`state = ipd_states,`
`year = ipd_year,`
`select(-NAME) dl_percs <- bind_cols(dl_percs, s_percs) }`
`if(length(d_calls$id > 0)){ d_percs <- get_acs(geography = "tract",`
`state = ipd_states,`
`year = ipd_year,`
`select(-NAME) dl_percs <- left_join(dl_percs, d_percs) }`
`if(length(dp_calls$id > 0)){ dp_percs <- get_acs(geography =`
`"tract",`
`output = "wide",`
`variables = dp_calls$api) %>% select(-NAME) dl_percs <-`
`left_join(dl_percs, dp_percs) } perc_s_calls$api <-`

```
str_replace(percs_calls$api, "PE", "") names(dl_percs) <-
str_replace(names(dl_percs), "PE", "E") names(dl_percs) <-
str_replace(names(dl_percs), "PM", "M") for(i in
1:length(percs_calls$id)){ names(dl_percs) <-
str_replace(names(dl_percs),
percs_calls$api[i],
percs_calls$id[i]) } dl_percs <- dl_percs %>% rename(GE0ID10 =
GE0ID)
```

5d. Format downloads

Subset `dl_counts` and `dl_percs` for DVRPC's nine-county region. Percentages should range from 0 to 100. {r `dl_counts_dl_percs`, message = FALSE} `dl_counts <- dl_counts %>% filter(str_sub(GE0ID10, 1, 5) %in% ipd_counties)`
`dl_percs <- dl_percs %>% filter(str_sub(GE0ID10, 1, 5) %in% ipd_counties)`

5d.i. Exception

Note that variable B02001_002 ("Estimate; Total: - White alone") was downloaded as the count for racial minority. Compute B02001_001 (Universe) – B02001_002 ("Estimate; Total: - White alone") and substitute for `RM_CE`. {r `perc_excp_1`, message = FALSE}
`dl_counts <- dl_counts %>% mutate(x = RM_UE - RM_CE) %>% select(-RM_CE) %>% rename(RM_CE = x)`

5d.ii. Exception

Before computing percentages and percentage MOEs, import the count MOE for the racial minority variable computed from variance replicates. If `rm_moe` exists, then this chunk will substitute the correct count MOE in `dl_counts`; if not, this chunk will do nothing. {r `perc_excp_2`, message = FALSE} `if(exists("rm_moe")){ dl_counts <- dl_counts %>% select(-RM_CM) %>% left_join(., rm_moe) %>% rename(RM_CM = RM_CntMOE) %>% mutate_at(vars(RM_CM), as.numeric) }`

5d.iii. Exception

Half-standard deviations serve as the classification bins for IPD scores, and including zero-population tracts affects computed standard deviation values. Start by removing the 11 census tracts with zero population. {r `perc_excp_3`} `slicer <- c("34005981802", "34005982200", "34021980000", "42017980000", "42045980300", "42045980000", "42045980200", "42091980100", "42091980000", "42091980200", "42091980300", "42101036901", "42101980001", "42101980002", "42101980003", "42101980300", "42101980701", "42101980702", "42101980800", "42101980100", "42101980200", "42101980400", "42101980500", "42101980600", "42101980901", "42101980902", "42101980903", "42101980904", "42101980905", "42101980906", "42101989100", "42101989200", "42101989300") dl_counts <- dl_counts %>% filter(!(GE0ID10 %in% slicer)) dl_percs <- dl_percs %>% filter(!(GE0ID10 %in% slicer))` Here are the first few lines of `dl_counts` and `dl_percs`. Notice the naming convention:

- UE = universe estimate
- UM = universe MOE
- CE = count estimate
- CM = count MOE
- PE = percentage estimate
- PM = percentage MOE

We use these strings to select columns, so consistency is key. `{r acs_preview}`
`head(dl_counts) head(dl_percs)`

6. ACS estimates calculations

For all nine indicators, this section computes:

- a. Percentages and percentage MOEs
- b. Percentile
- c. IPD score and classification
- d. Composite IPD score

Split `dl_counts` into a list named `comp` for processing and arrange column names in alphabetical order. The name of the list, `comp`, is a nod to the “component parts” of `dl_counts`. The structure of `comp` is similar to a four-tab Excel spreadsheet: for example, `comp` is the name of the `.xlsx` file, `uni_est` is a tab for universe estimates, and `uni_est` has nine columns and 1,368 rows, where the column is the IPD indicator and the row is the census tract observation.

The order of columns is important because processing is based on vector position. We want to make sure that the first column of every tab corresponds to the Disabled indicator, the second to Ethnic Minority, et cetera. `{r comp}`

```
comp <- list() comp$uni_est <- dl_counts %>% select(ends_with("UE")) %>% select(sort(current_vars()))
comp$uni_moe <- dl_counts %>% select(ends_with("UM")) %>% select(sort(current_vars()))
comp$count_est <- dl_counts %>% select(ends_with("CE")) %>% select(sort(current_vars()))
comp$count_moe <- dl_counts %>% select(ends_with("CM")) %>% select(sort(current_vars()))
```

6a. Percentages and percentage MOEs

6a.i. Calculation

MOEs of the percentage values are obtained using the `tidycensus` function `moe_prop`. This chunk mentions `r` and `c` several times: continuing the spreadsheet analogy, think of `r` as the row number and `c` as the column number for a given spreadsheet tab. `{r perc}`

```
pct_matrix <- NULL pct_moe_matrix <- NULL for (c in 1:length(comp$uni_est)){
  pct <- unlist(comp$count_est[,c] / comp$uni_est[,c])
  pct_matrix <- cbind(pct_matrix, pct)
  moe <- NULL
  for (r in 1:length(comp$uni_est$LI_UE)){
    moe_indiv <-
```

```
as.numeric(moe_prop(comp$count_est[r,c],
comp$uni_est[r,c],
comp$count_moe[r,c],
comp$uni_moe[r,c])) moe <- append(moe, moe_indiv) }
pct_moe_matrix <- cbind(pct_moe_matrix, moe) }
```

6a.ii. Result

pct and pct_moe stores the percentages and associated MOEs for the nine indicator variables. Results are rounded to the tenths place and range from 0 to 100. {r perc_res, warning = FALSE} pct <- as_tibble(pct_matrix) %>% mutate_all(funs(. * 100)) %>% mutate_all(round, 1) names(pct) <- str_replace(names(comp\$uni_est), "_UE", "_PctEst") pct_moe <- as_tibble(pct_moe_matrix) %>% mutate_all(funs(. * 100)) %>% mutate_all(round, 1) names(pct_moe) <- str_replace(names(comp\$uni_est), "_UE", "_PctMOE")

6a.iii. Exception

If the percentage MOE equals 0, then overwrite it to equal 0.1. This should be a rare occurrence with survey data at the census tract level. {r perc_excp_4} pct_moe <- pct_moe %>% replace(., . == 0, 0.1)

6a.iv. Exception

Substitute percentages and associated MOEs when available. This applies to the older adults, female, limited English proficiency, and disabled variables. {r perc_excp_5} pct <- pct %>% mutate(D_PctEst = dl_percs\$D_PE, OA_PctEst = dl_percs\$OA_PE, LEP_PctEst = dl_percs\$LEP_PE, F_PctEst = dl_percs\$F_PE) pct_moe <- pct_moe %>% mutate(D_PctMOE = dl_percs\$D_PM, OA_PctMOE = dl_percs\$OA_PM, LEP_PctMOE = dl_percs\$LEP_PM, F_PctMOE = dl_percs\$F_PM) Here are the first few lines of pct and pct_moe: {r pct_preview} head(pct) head(pct_moe)

6b. Percentile

6b.i. Calculation

Add percentiles (an additional “spreadsheet tab”) to comp, making sure to first sort column names alphabetically. Compute the empirical cumulative distribution function for each of the nine indicator variables. The ECDF can range from 0 to 1, where 1 indicates the largest observed percentage. {r percentile} comp\$pct_est <- pct %>% select(sort(current_vars())) percentile_matrix <- NULL for (c in 1:length(comp\$uni_est)){ p <- unlist(comp\$pct_est[,c]) rank <- ecdf(p)(p) percentile_matrix <- cbind(percentile_matrix, rank) }

6b.ii. Result

percentile stores the percentile for the nine indicator variables. Results are rounded to the hundredths place. {r percentile_res, warning = FALSE} percentile <- as_tibble(percentile_matrix) %>% mutate_all(round, 2)
names(percentile) <- str_replace(names(comp\$uni_est), "_UE", "_Pctile") Here are the first few lines of percentile: {r percentile_preview}
head(percentile)

6c. IPD score and classification

Each observation is assigned an IPD score for each indicator. The IPD score for an individual indicator can range from 0 to 4, which corresponds to the following classification and bin breaks:

```
{r table4, tidy=FALSE} library(tidyverse) data4 <- tribble( ~"IPD  
Score", ~"IPD Classification", ~"Standard Deviations", "0", "Well  
Below Average", "x $< -1.5 \cdot stdev$", "1", "Below Average", "$-1.5  
\cdot stdev \leq$ x $<-0.5 \cdot stdev$", "2", "Average", "$-0.5 \cdot  
stdev \leq$ x $<0.5 \cdot stdev$", "3", "Above Average", "$0.5 \cdot  
stdev \leq$ x $<1.5 \cdot stdev$", "4", "Well Above Average", "x $\geq  
1.5 \cdot stdev$" ) knitr::kable((data4), booktabs = TRUE, caption =  
'')
```

6c.i. Calculation

The function `st_dev_breaks` is called to compute the bin breaks for each indicator. These breaks determine the IPD score stored in `score`. Note that we divide *rounded* `PctEst` columns by *unrounded* half-standard deviation breaks to compute the score. `class` is a textual explanation of the IPD score. {r score_class} score_matrix <- NULL
class_matrix <- NULL for (c in 1:length(comp\$uni_est)){ p <-
unlist(comp\$pct_est[,c]) breaks <- st_dev_breaks(p, 5, na.rm = TRUE)
score <- case_when(p < breaks[2] ~ 0, p >= breaks[2] & p < breaks[3] ~ 1,
breaks[3] & p < breaks[4] ~ 2, p >= breaks[4] & p < breaks[5] ~ 3,
p >= breaks[5] ~ 4) class <- case_when(score
== 0 ~ "Well Below Average", score == 1 ~ "Below
Average", score == 2 ~ "Average",
score == 3 ~ "Above Average", score == 4 ~ "Well
Above Average") score_matrix <- cbind(score_matrix, score)
class_matrix <- cbind(class_matrix, class) }

6c.ii. Result

`score` and `class` store the IPD scores and associated descriptions for the nine indicator variables. {r score_class_res, warning = FALSE} score <- as_tibble(score_matrix) names(score) <-
str_replace(names(comp\$uni_est), "_UE", "_Score") class <- as_tibble(class_matrix) names(class) <-

`str_replace(names(comp$uni_est), "_UE", "_Class")` Here are the first few lines of score and class: `{r score_preview} head(score) head(class)`

6d. Composite IPD score

6d.i. Calculation

Sum the IPD scores for the nine indicator variables to determine the composite IPD score. `{r ipd_score} score <- score %>% mutate(IPD_Score = rowSums(.))`

6d.ii. Result

Here are the first few records of the composite IPD score: `{r ipd_score_preview} head(score$IPD_Score)`

7. ACS estimates cleaning

There is a specific output format for `ipd.csv`, including column names, column order, flags for missing data, and census tracts with insufficient data. This section ensures conformity with the output formatting.

Merge the percentage estimates, percentage MOEs, percentile, score, and class data frames into a single data frame called `ipd`. `{r merge} ipd <- bind_cols(dl_counts, pct) %>% bind_cols(., pct_moe) %>% bind_cols(., percentile) %>% bind_cols(., score) %>% bind_cols(., class)` Rename columns. `{r rename} names(ipd) <- str_replace(names(ipd), "_CE", "_CntEst") names(ipd) <- str_replace(names(ipd), "_CM", "_CntMOE") ipd <- ipd %>% mutate(STATEFP10 = str_sub(GEOID10, 1, 2), COUNTYFP10 = str_sub(GEOID10, 3, 5), NAME10 = str_sub(GEOID10, 6, 11), U_TPopEst = F_UE, U_TPopMOE = F_UM, U_Pop5Est = LEP_UE, U_Pop5MOE = LEP_UM, U_PPovEst = LI_UE, U_PPovMOE = LI_UM, U_PNICEst = D_UE, U_PNICMOE = D_UM) %>% select(-ends_with("UE"), -ends_with("UM"))` Reorder columns, with `GEOID` and `FIPS` codes first, the following variables in alphabetical order, and the total IPD score and universes at the end. `{r reorder} ipd <- ipd %>% select(GEOID10, STATEFP10, COUNTYFP10, NAME10, sort(current_vars())) %>% select(move_last(., c("IPD_Score", "U_TPopEst", "U_TPopMOE", "U_Pop5Est", "U_Pop5MOE", "U_PPovEst", "U_PPovMOE", "U_PNICEst", "U_PNICMOE")))` At the beginning of processing, we removed 11 census tracts from processing because their populations were equal to zero. Tack these back on to the dataset. `{r tack} slicer <- enframe(slicer, name = NULL, value = "GEOID10") ipd <- plyr::rbind.fill(ipd, slicer)` Replace NA values with `NoData` if character and `-99999` if numeric. `{r replace} ipd <- ipd %>% mutate_if(is.character, funs(ifelse(is.na(.), "NoData", .))) %>% mutate_if(is.numeric, funs(ifelse(is.na(.), -99999, .)))`

8. Summary Tables

This section generates a handful of other deliverables, including:

- a. Counts by indicator
- b. Breaks by indicator
- c. Summary by indicator
- d. County means by indicator

Replace -99999 with NA for numeric columns to avoid distorting summary statistics. {r summary_prep} ipd_summary <- ipd ipd_summary[ipd_summary == -99999] <- NA

8a. Counts by indicator

The number of census tracts that fall in each bin. Count census tracts by indicator and bin. Reorder factor levels so that “Well Below Average” appears before “Below Average,” and the like. {r summary_counts, message = FALSE, warning = FALSE} counts <- ipd_summary %>% select(ends_with("Class")) export_counts <- apply(counts, 2, function(i) plyr::count(i)) for(i in 1:length(export_counts)){ export_counts[[i]]\$var <- names(export_counts)[i] } export_counts <- map_dfr(export_counts, `[`, c("var", "x", "freq")) colnames(export_counts) <- c("Variable", "Classification", "Count") export_counts\$Classification <- factor(export_counts\$Classification, levels = c("Well Below Average", "Below Average", "Average", "Above Average", "NoData")) export_counts <- arrange(export_counts, Variable, Classification) export_counts <- export_counts %>% spread(Classification, Count) %>% mutate_all(funs(replace_na(., 0))) %>% mutate(TOTAL = rowSums(.[2:7], na.rm = TRUE))

8b. Breaks by indicator

The bin breaks for each indicator. Apply the st_dev_breaks function to all percentage values and export results. {r summary_breaks} breaks <- ipd_summary %>% select(ends_with("PctEst")) export_breaks <- round(mapply(st_dev_breaks, x = breaks, i = 5, na.rm = TRUE), digits = 3) export_breaks <- as_tibble(export_breaks) %>% mutate(Class = c("Min", "1", "2", "3", "4", "Max")) %>% select(Class, current_vars())

8c. Summary by indicator

Summary statistics of each indicator. Round results to two decimal places. {r summary_summary} pcts <- ipd_summary %>% select(ends_with("PctEst"))


```
summary_data <- apply(pcts, 2, description) export_summary <-
as_tibble(summary_data) %>% mutate_all(round, 2) %>%
mutate(Statistic = c("Minimum", "Median", "Mean", "SD", "Half-SD",
"Maximum")) %>% select(Statistic, current_vars())
```

8d. County means by indicator

Population-weighted means by county and indicator. For the most accurate percentage values, aggregate all counts back to the county level and compute percentages. In the export file, counties are referred to by the five-digit character supplied by the user to `ipd_counties`. {r summary_county, warning = FALSE, message = FALSE}

```
export_means <- dl_counts %>% select(GEOID10, ends_with("UE"),
ends_with("CE")) %>% select(GEOID10, sort(current_vars())) %>%
mutate(County = str_sub(GEOID10, 1, 5)) %>% select(-GEOID10) %>%
group_by(County) %>% summarize(D_PctEst = sum(D_CE) / sum(D_UE),
EM_PctEst = sum(EM_CE) / sum(EM_UE), F_PctEst =
sum(F_CE) / sum(F_UE), FB_PctEst = sum(FB_CE) /
sum(FB_UE), LEP_PctEst = sum(LEP_CE) / sum(LEP_UE),
LI_PctEst = sum(LI_CE) / sum(LI_UE), OA_PctEst =
sum(OA_CE) / sum(OA_UE), RM_PctEst = sum(RM_CE) /
sum(RM_UE), Y_PctEst = sum(Y_CE) / sum(Y_UE)) %>%
mutate_if(is.numeric, funs(. * 100)) %>% mutate_if(is.numeric,
round, 1)
```

9. Export

9a. Append to TIGER/LINE file

Using the arguments supplied in `ipd_county`, download the relevant census tracts and append `ipd` to them. Uncommenting `cb = TRUE` will greatly speed processing time by downloading generalized tract boundary shapefiles instead of detailed ones. {r

```
shapefile, message = FALSE, warning = FALSE} options(tigris_use_cache
= TRUE, tigris_class = "sf") st <- str_sub(ipd_counties, 1, 2) cty <-
str_sub(ipd_counties, 3, 5) trct <- map2(st, cty, ~{tracts(state = .x,
county = .y,
#cb = TRUE,
year = ipd_year)}) %>% rbind_tigris() %>% st_transform(., 26918)
%>% select(GEOID) %>% left_join(., ipd, by = c("GEOID" =
"GEOID10")) %>% rename(GEOID10 = GEOID)
```

9b. Export files

Results are saved in `outputs`. {r happy_trails, message = FALSE, warning = FALSE}

```
st_write(trct, here("outputs", "ipd.shp"), delete_dsn = TRUE,
quiet = TRUE) write_csv(ipd, here("outputs", "ipd.csv"))
write_csv(export_counts, here("outputs", "counts_by_indicator.csv"))
write_csv(export_breaks, here("outputs", "breaks_by_indicator.csv"))
write_csv(export_summary, here("outputs", "summary_by_indicator.csv"))
write_csv(export_means, here("outputs", "mean_by_county.csv"))
```